

## 1 Introduction

We propose a new model for dimensionality reduction, the PCA tree, which works like a regular autoencoder, having explicit projection and reconstruction mappings. The projection is effected by a sparse oblique tree, having hard, hyperplane splits using few features and linear leaves. The reconstruction mapping is a set of local linear mappings.

Work partially supported by NSF award IIS-2007147.

## 3 Optimization and time complexity

Our objective function is the regular reconstruction error of an autoencoder  $\mathbf{T}: \mathbb{R}^D \rightarrow \mathbb{R}^D$  with an  $\ell_1$  regularization term of hyperparameter  $\lambda \geq 0$  on a training set  $\{\mathbf{x}_n\}_{n=1}^N \subset \mathbb{R}^D$ :

$$E(\Theta) = \sum_{n=1}^N \|\mathbf{x}_n - \mathbf{T}(\mathbf{x}_n; \Theta)\|_2^2 + \lambda \sum_{i \in D} \|\mathbf{w}_i\|_1 \quad \text{s.t.} \quad \mathbf{U}_j^T \mathbf{U}_j = \mathbf{I}, \forall j \in \mathcal{L}.$$

We use a variation of the Tree Alternating Optimization (TAO) approach to train the model. It repeatedly updates the nodes in turn: at a leaf it solves a PCA, and at a decision node it solves an  $\ell_1$ -regularized, logistic regression problem. The training complexity is at most linear in  $N$  and quadratic in  $D$ , just like in regular PCA. As a useful summary, *the rough cost is  $\Theta(ND^2)$  for shallow trees and  $\mathcal{O}(ND)$  for deep trees*—which is asymptotically faster than PCA!

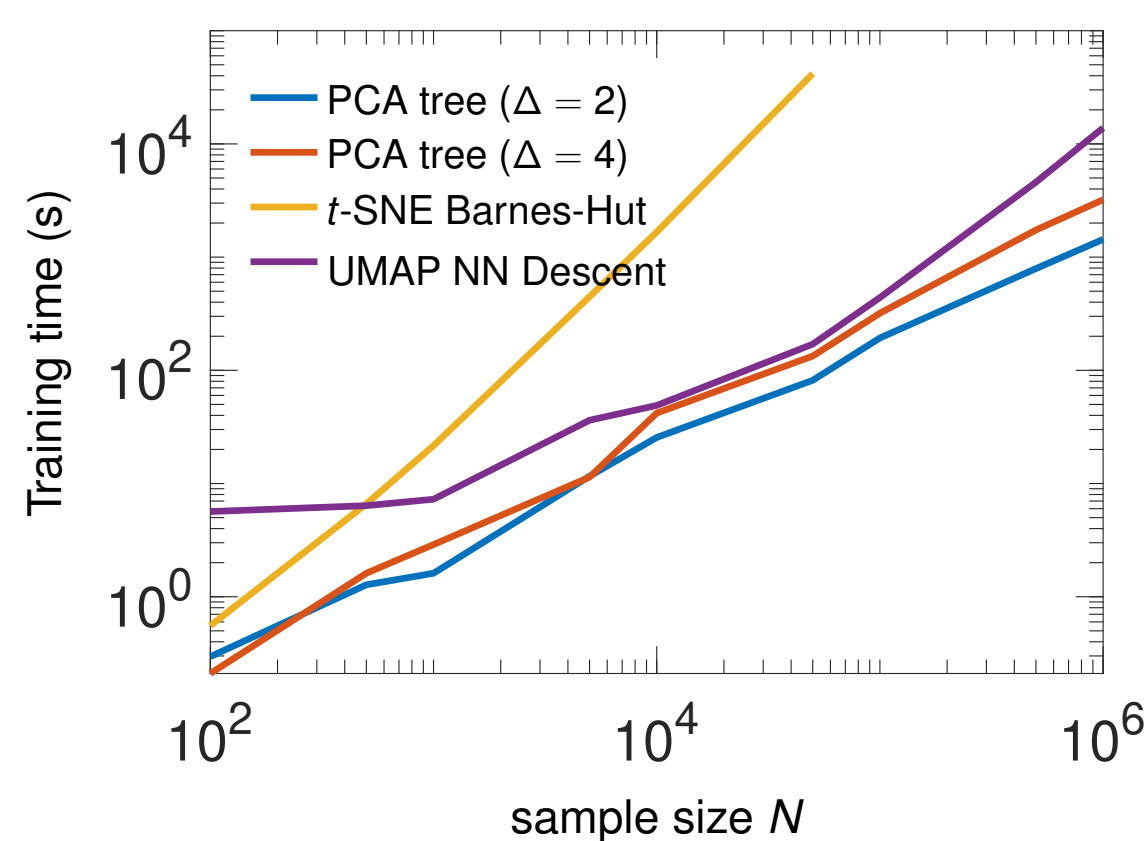


Figure: Training time per iteration on Infinite MNIST for PCA trees with different  $N$ ,  $\Delta$ , and for  $t$ -SNE and UMAP.

## 5 Fashion MNIST visualizations

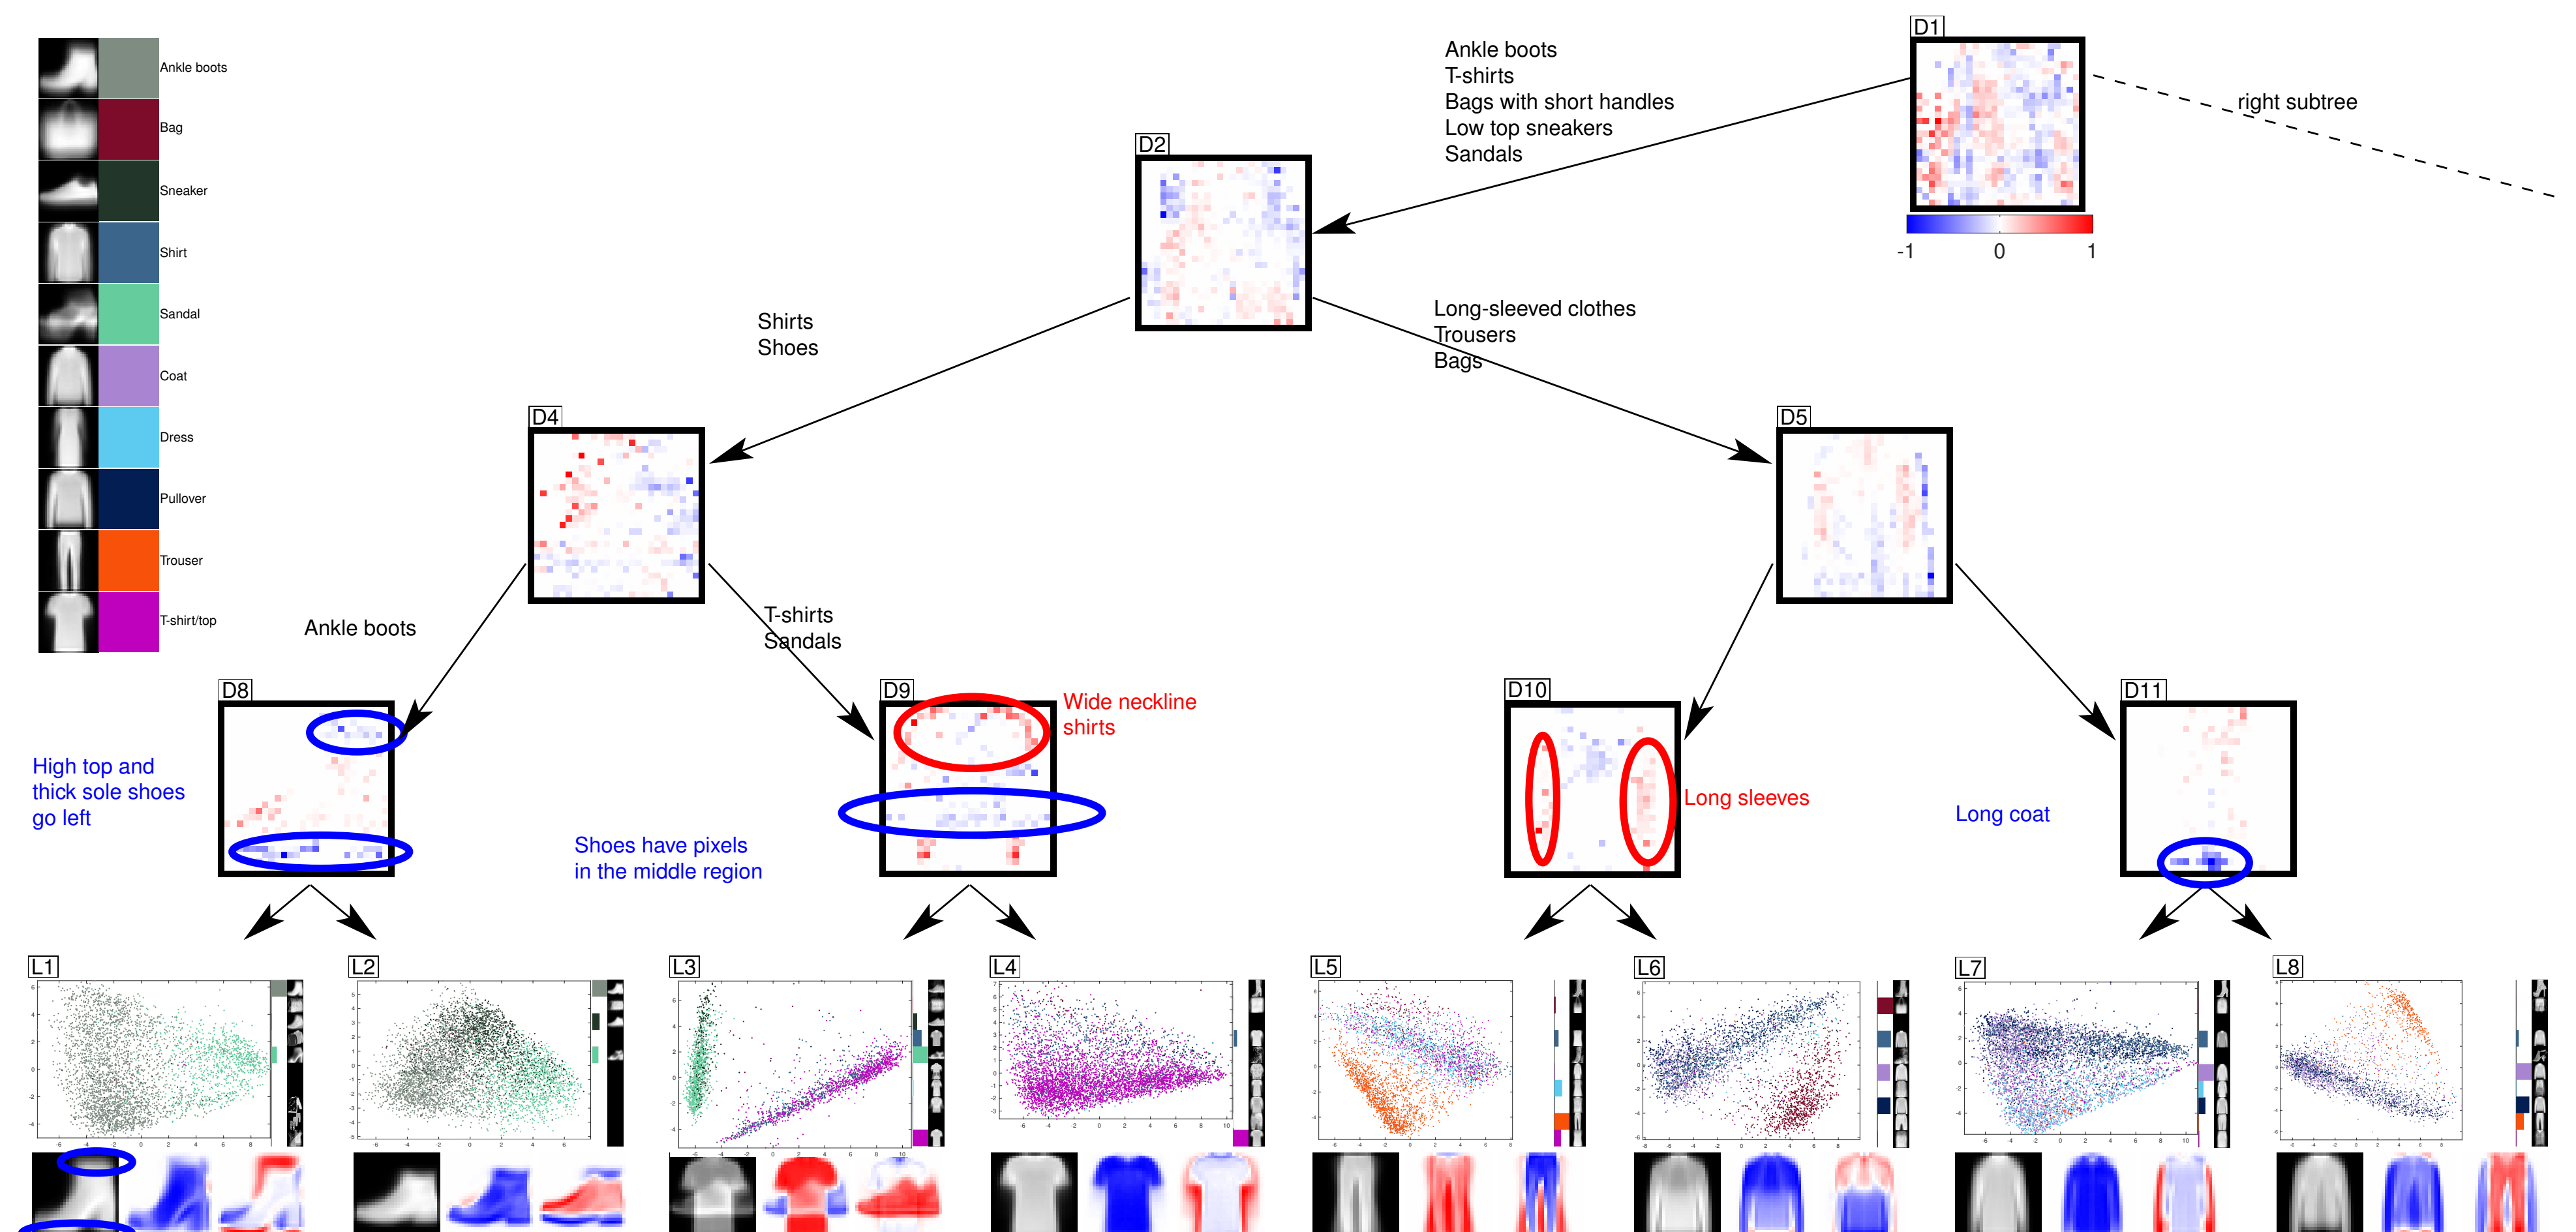
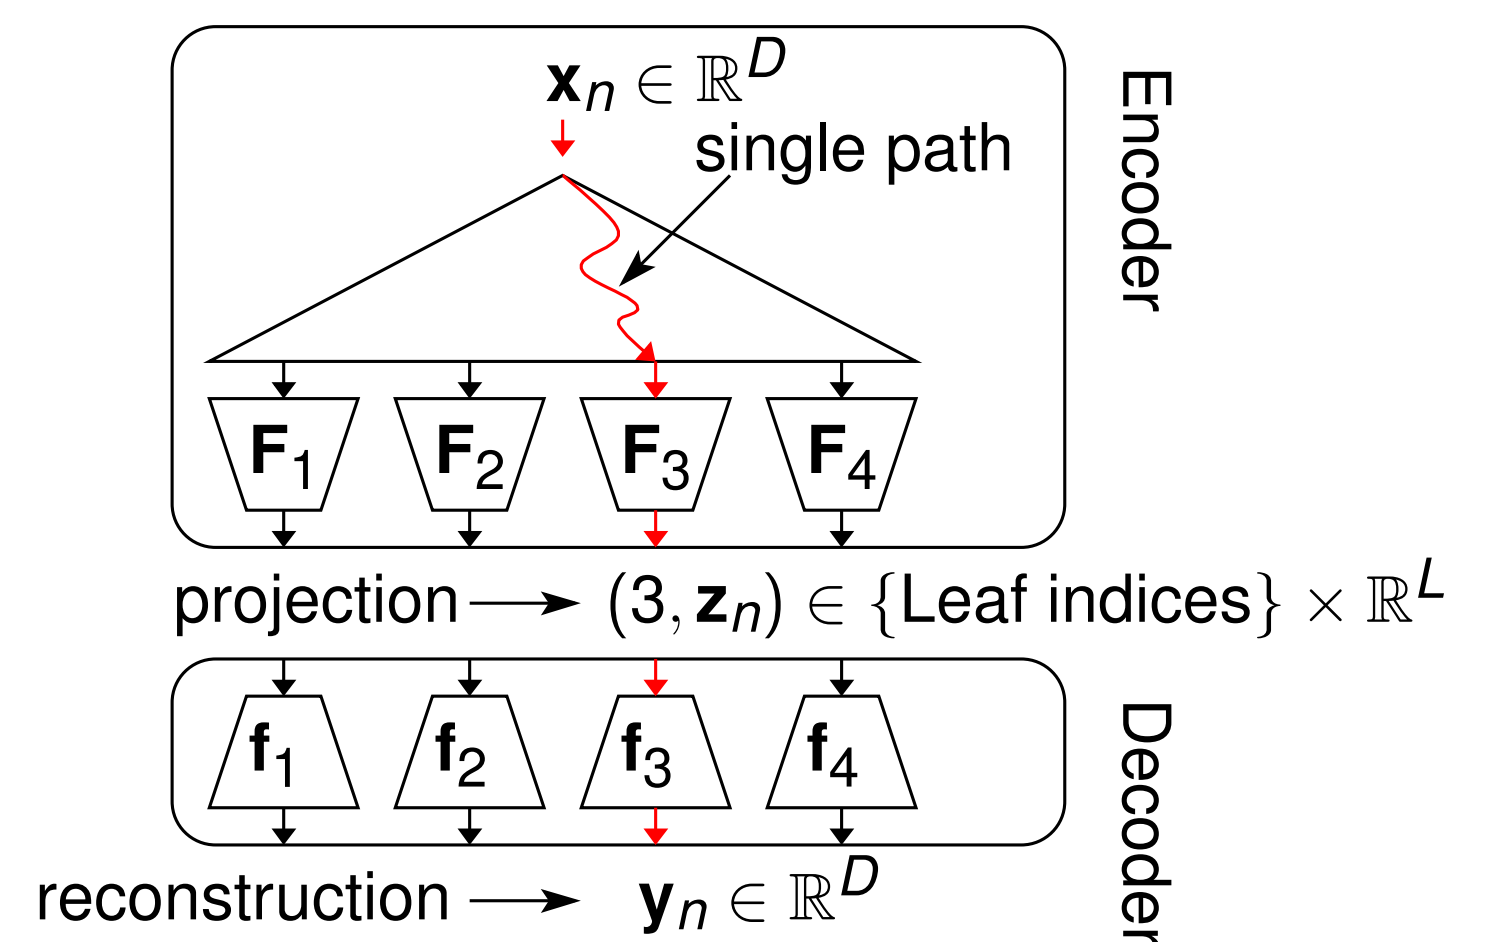


Figure: PCA tree trained on Fashion MNIST. The decision nodes' weight vectors (and the leaves' PCs  $\mathbf{U}_j$ ) are shown as  $28 \times 28$  images, with negative/zero/positive values colored blue/white/red. Each leaf shows a 2D PCA scatterplot of its RS, and below it the mean  $\mu_j$  (grayscale) image and the 2 PCs  $\mathbf{U}_j$  (in color). To the right of the scatterplot, a bar chart displays class proportions and class means. The legend (top left) shows each class's grayscale mean and color (for the scatterplots).

## 2 Tree autoencoder

The **encoder** is given by a tree mapping  $\mathbf{T}^e(\mathbf{x}; \Theta): \mathbb{R}^D \rightarrow \mathcal{L} \times \mathbb{R}^L$  where the predictor for leaf  $j$  has the form of a linear mapping  $\mathbf{F}_j(\mathbf{x}; \mathbf{U}_j, \mu_j) = \mathbf{U}_j^T(\mathbf{x} - \mu_j)$ , where  $\mathbf{U}_j \in \mathbb{R}^{D \times L}$  is an orthogonal matrix and  $\mu_j \in \mathbb{R}^D$ . The encoder parameters are  $\Theta = \{\mathbf{w}_i, w_{i0}\}_{i \in D} \cup \{\mathbf{U}_j, \mu_j\}_{j \in \mathcal{L}}$ . Thus, the encoder maps an input instance  $\mathbf{x} \in \mathbb{R}^D$  to a leaf index  $j \in \mathcal{L}$  and an  $L$ -dimensional real vector  $\mathbf{z} = \mathbf{U}_j^T(\mathbf{x} - \mu_j)$ , which at an optimum will be the PCA projection in that leaf. This means that the PCA tree does not have a common latent space of dimension  $L$  where all instances are projected. Instead, it has one separate  $L$ -dimensional PCA space per leaf.

The **decoder** maps a leaf index  $j$  and  $L$ -dimensional vector  $\mathbf{z}$  (in  $\mathcal{L} \times \mathbb{R}^L$ ) to a vector in  $\mathbb{R}^D$ . It consists of a set of linear mappings of the form  $\mathbf{f}_j(\mathbf{z}; \mathbf{U}_j, \mu_j) = \mathbf{U}_j \mathbf{z} + \mu_j$  for  $j \in \mathcal{L}$ .



## 4 Advantages

The PCA tree provides significant, complementary advantages over previous methods:

1. It optimizes the reconstruction error, which has a clear meaning
2. It does not require a neighborhood graph (and perplexity parameter, etc.), which is tricky to estimate so it captures manifold structure, and computationally very costly
3. It is highly interpretable and extracts a wealth of information from complex datasets.
4. It defines nonlinear out-of-sample mappings
5. PCA map shows clusters, these are real—in contrast with  $t$ -SNE's tendency to create false clusters
6. The loss function is really a self-supervised regression problem. This makes it possible to use cross-validation to determine the hyperparameters